

Nathan Kim

(+1)-650-505-4069 | nathangk@stanford.edu | [linkedin.com/in/nathangk](https://www.linkedin.com/in/nathangk) | [nathankim7.github.io](https://github.com/nathankim7)

EDUCATION

Stanford University

BS, Symbolic Systems

Stanford, CA

Class of 2024

Key Coursework: NLP w. Deep Learning (Best Project), Machine Learning (Best Project), Deep Generative Models

RESEARCH

Holistic Evaluation of Language Models

Transactions on Machine Learning Research (TMLR)

Authored with Stanford CRFM

This paper introduces HELM, a **living benchmark** for foundation models designed to capture the full space of LLM use cases in a principled manner. **Individually**, I implement the International Corpus of English (ICE) as an evaluation on language modelling fairness across English dialects of varying prestige, and argue for its importance in an assessment of bias expressed by LLMs. I also discuss the results of our suite of linguistic knowledge evaluations more generally.

GLARE: Infilling Language Models for Textual Adversarial Attacks

Eval4NLP @ ACL-IJCNLP 2022

Nathan Kim, Ryan Chi*, Patrick Liu, Zander Lack, Ethan Chi*

Textual adversarial examples can expose serious security vulnerabilities in NLP models by illustrating where small perturbations in input texts can drastically alter their predictions. We find that the GPT-2 adapted to perform non-causal infilling (Donahue et al. 2020) outperforms existing adversarial example generation methods in fluency, semantic preservation, and attack success rates.

WORK EXPERIENCE

Unify GTM

July 2024—

Software Engineer

- **Shipped** features start-to-finish including usage upsells and upgraded **LM-powered copy generation** as 6th engineering hire.
- **Spearheaded** optimization of query-caching and stylesheet loading among other foundational frontend technical contributions

Portalform (now Fiber) (YCombinator W23)

January 2023—March 2023

Engineering Intern

- **Bootstrapped** a next-generation software layer for SaaS to automate 3rd-party client data ingestion.
- **Designed** a reader service to sync client databases with Amazon + Shopify. Built on Prisma, PostgreSQL & Express.

Canal

June 2022—September 2022

Frontend Engineering Intern

- Contributed **18.5% of frontend code** written at Canal during internship period, including **redesigned invite interface**
- **Initiated/designed major refactor** of Proposal screens and provided **key input on data model** for V2 Proposal system

PROJECTS & VOLUNTEER WORK

Causal Alignment for Controlled Text Generation

October 2023—June 2024

Lead Researcher, advised by Zhengxuan Wu, Christopher Potts

- Thesis research on developing prompt-free one-click controls for natural language generation (NLG) with LLMs
- **Expands** on theory of IIT (Geiger et al. 2022) to control for causal variables (sentiment, phrase structure) in generated sequences of infinite length
- Matched accuracy + fluency of existing CTG methods with **zero** new trained parameters
- **Initiated/designed** implementation of all methods, to appear in open-source ML library [Pyvene](#)

AWARDS

- Gold Medal Division, 2020 Asia Pacific Linguistics Olympiad
- Silver Medalist, 2019 International Linguistics Olympiad
- Student Honor Roll, 2018 Canadian Computing Competition

SKILLS

- Pytorch, Pandas, R, HTML/CSS. React.js, Flask, Node.js, Java, React Native, C++, Rust